



Stage 2 — ConvNeXt-Tiny, v1 vs v2


Validation macro F1 (checkpoint selection metric)



Accuracy — the train/val gap is overfitting



Loss



Learning-rate schedule

